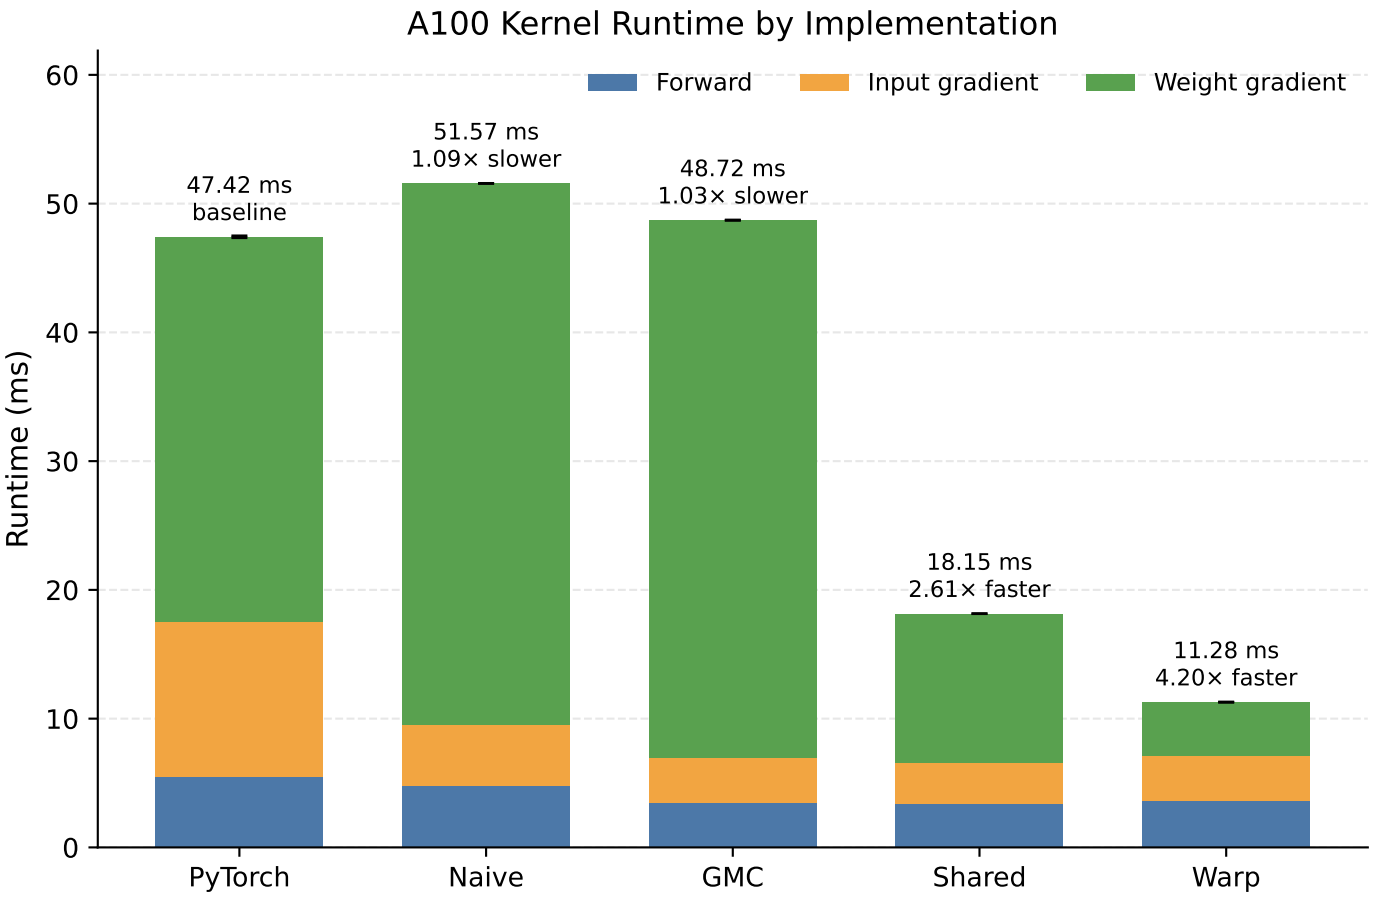
A100 Kernel Runtime by Implementation
Forward
Input gradient
Weight gradient
Runtime (ms)
0
10
20
30
40
50
60
47.42 ms
baseline
51.57 ms
1.09× slower
48.72 ms
1.03× slower
18.15 ms
2.61× faster
11.28 ms
4.20× faster
PyTorch
Naive
GMC
Shared
Warp