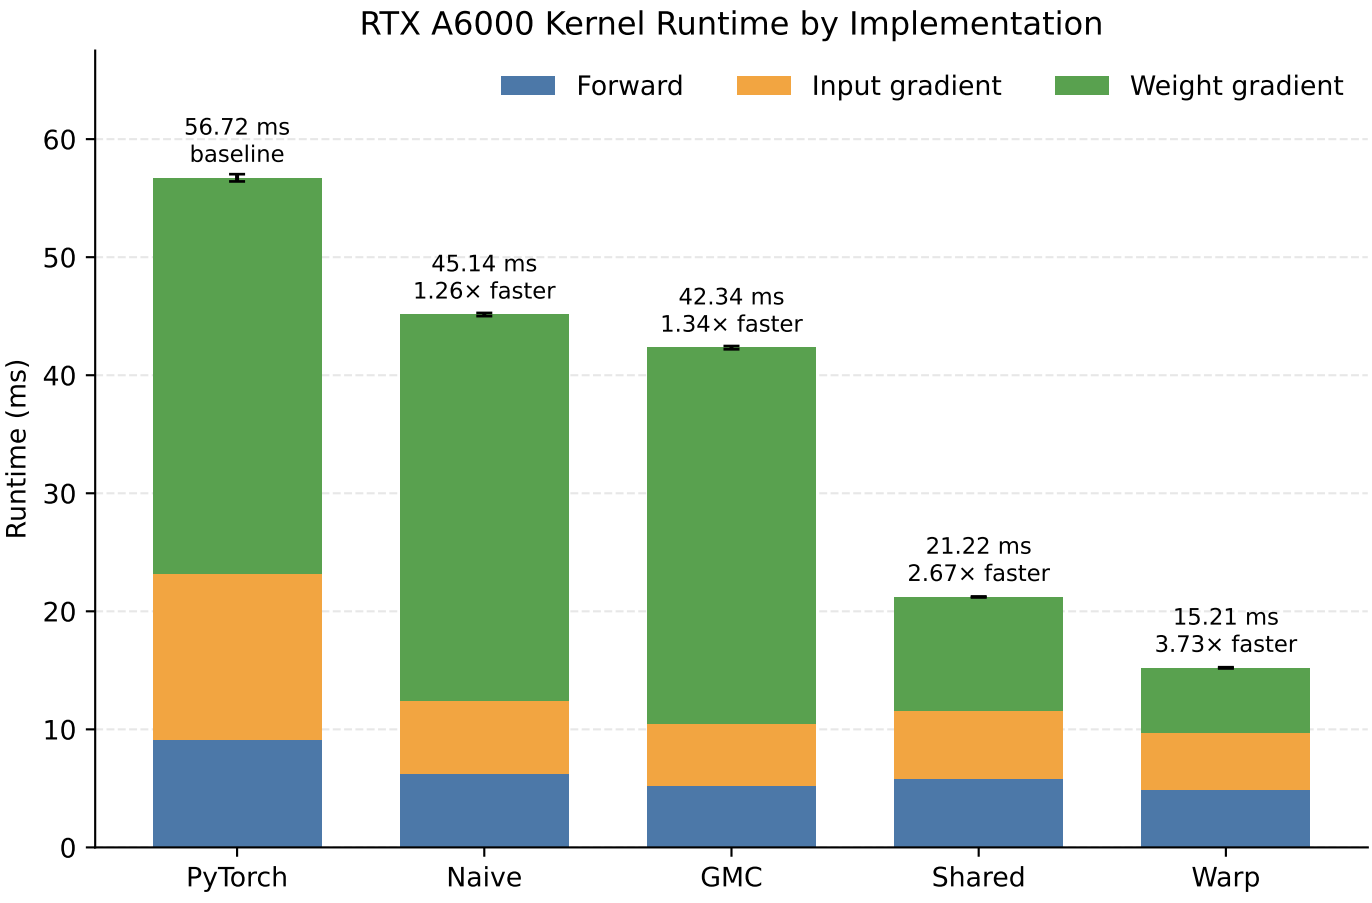
RTX A6000 Kernel Runtime by Implementation
Forward
Input gradient
Weight gradient
Runtime (ms)
0
10
20
30
40
50
60
56.72 ms
baseline
45.14 ms
1.26× faster
42.34 ms
1.34× faster
21.22 ms
2.67× faster
15.21 ms
3.73× faster
PyTorch
Naive
GMC
Shared
Warp